



ECE408 Fall 2022


# Applied Parallel Programming

## Lecture 17 Parallel Scan Part-2

# Objective

- To master parallel scan (prefix sum) algorithms
  - Work-efficiency vs. latency
  - Brent-Kung Tree Algorithm
  - Hierarchical algorithms

# A Kogge-Stone Parallel Scan Algorithm




# Improving Efficiency

- A common parallel algorithm pattern:

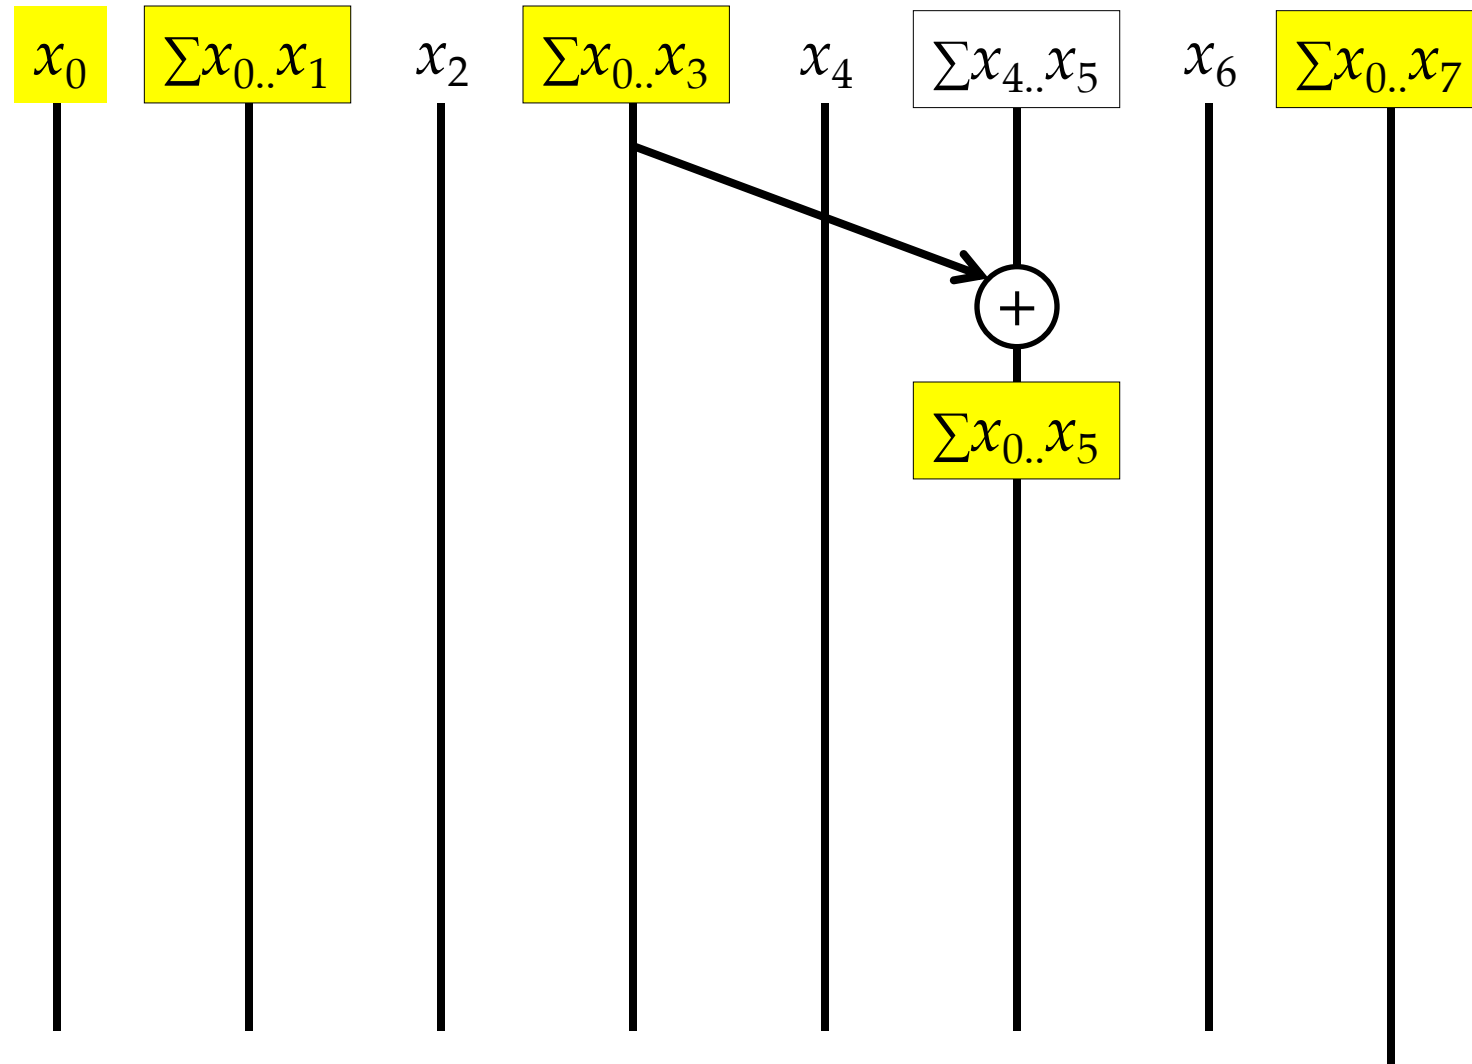
## *Balanced Trees*

- Build a balanced binary tree on the input data and sweep it to and from the root
  - Tree is not an actual data structure, but a conceptual pattern
- 
- For scan:
    - Traverse down from leaves to root building partial sums at internal nodes in the tree
      - Root holds sum of all leaves
    - Traverse back up the tree building the scan from the partial sums

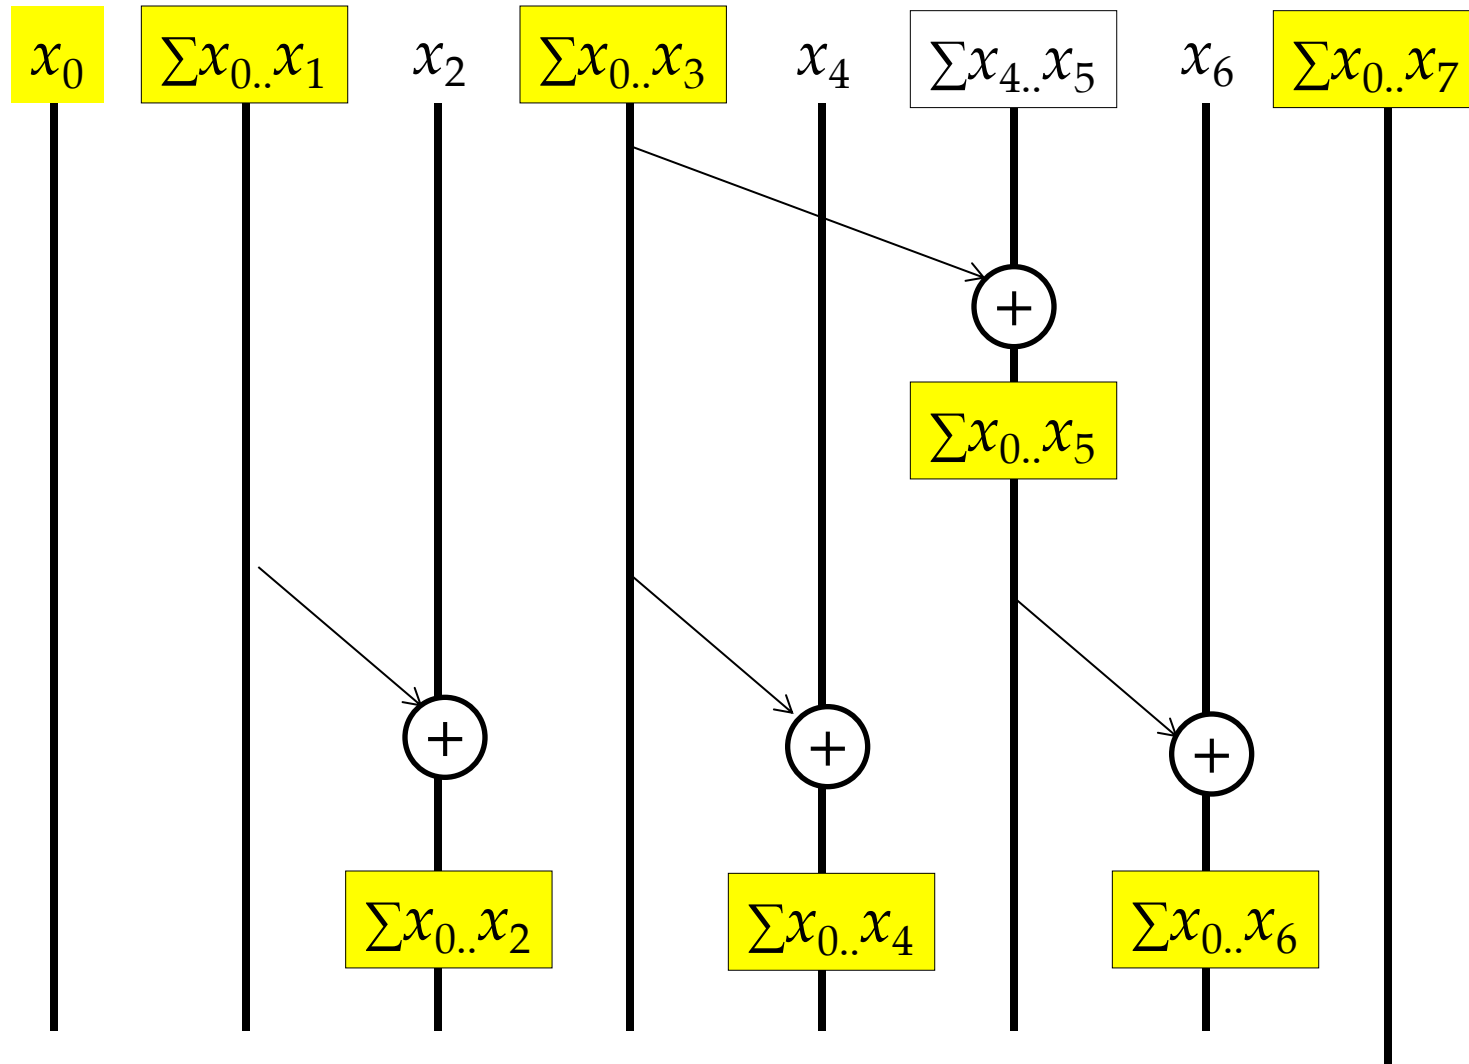
# Brent-Kung Parallel Scan Step




# Post-Scan Step



# Post Scan Step



# Putting it Together (Data View)






# Scan Step Kernel Code

```
// float T[2*BLOCK_SIZE] is in shared memory
// for previous slide, BLOCK_SIZE is 8
int stride = 1;
while(stride < 2*BLOCK_SIZE) {
    __syncthreads();
    int index = (threadIdx.x+1)*stride*2 - 1;
    if(index < 2*BLOCK_SIZE && (index-stride) >= 0)
        T[index] += T[index-stride];
    stride = stride*2;
}
```

```
// In our example,
// threadIdx.x+1    = 1, 2, 3, 4, 5, 6, 7, 8
// stride = 1, index = 1, 3, 5, 7, 9, 11, 13, 15
```

# Putting it Together



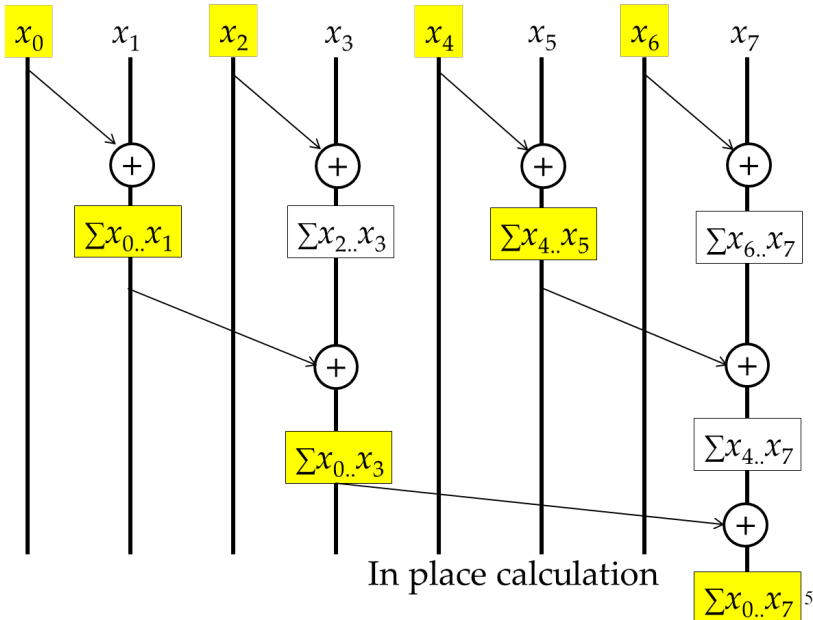
# Post Scan Step

```
int stride = BLOCK_SIZE/2;
while(stride > 0) {
    __syncthreads();
    int index = (threadIdx.x+1)*stride*2 - 1;
    if ((index+stride) < 2*BLOCK_SIZE)
        T[index+stride] += T[index];
    stride = stride / 2;
}

// In our example,
// BLOCK_SIZE=8 stride=4, 2, 1
// for first iteration, active thread = 0 index = 7, +stride = 11
```

# Work Analysis

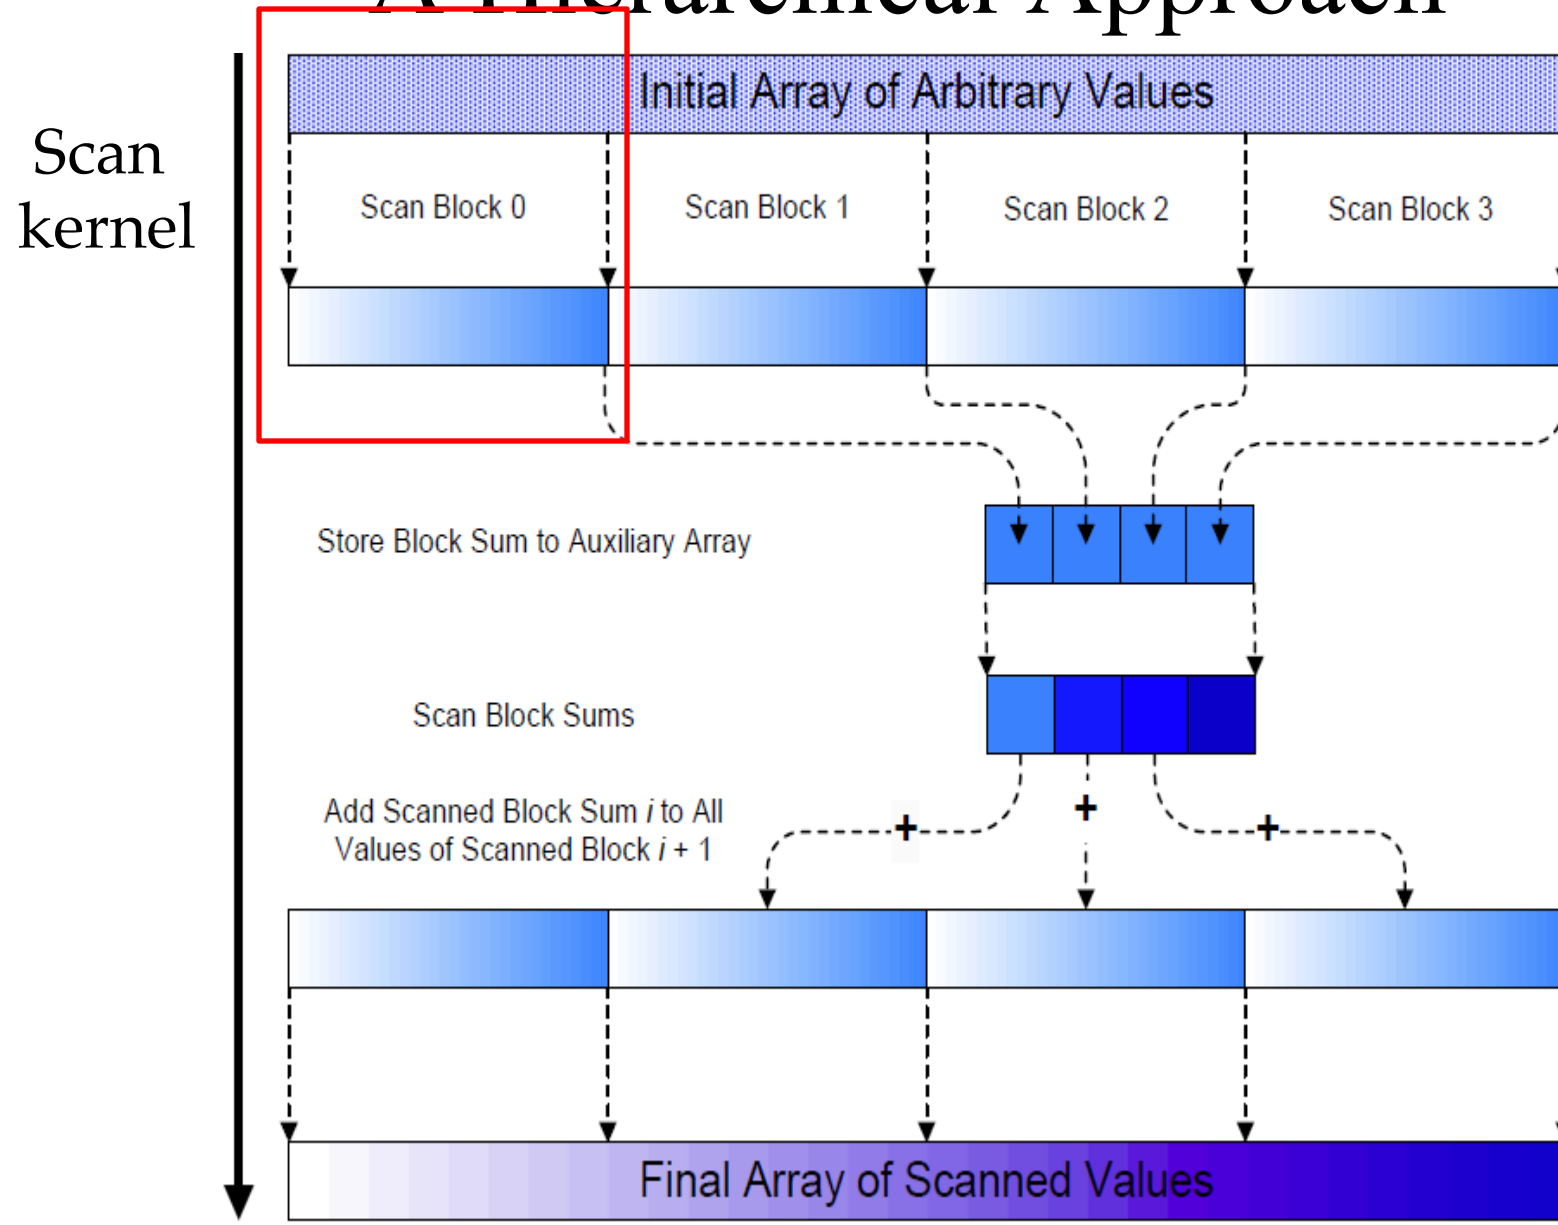
- The parallel Scan executes  $2 * \log(n)$  parallel iterations
  - $\log(n)$  in reduction and  $\log(n)$  in post scan
  - The iterations do  $n/2, n/4, \dots, 1, (2-1), \dots, (n/4-1), (n/2-1)$  useful adds
  - In our example,  $n = 16$ , the number of useful adds is  $16/2 + 16/4 + 16/8 + 16/16 + (16/8-1) + (16/4-1) + (16/2-1)$
  - Total adds:  $(n-1) + (n-2) - (\log(n) - 1) = 2*(n-1) - \log(n) \rightarrow O(n)$  work
- The total number of adds is no more than twice of that done in the efficient sequential algorithm
  - The benefit of parallelism can easily overcome the 2X work when there is sufficient hardware



- 13

# Overall Flow of Complete Scan

## A Hierarchical Approach



# Using Global Memory Contents in CUDA

- Data in registers and shared memory of one thread block are not visible to other blocks
- To make data visible, the data has to be written into global memory
- However, any data written to the global memory are not visible until a memory fence. This is typically done by terminating the kernel execution
- Launch another kernel to continue the execution. The global memory writes done by the terminated kernels are visible to all thread blocks.

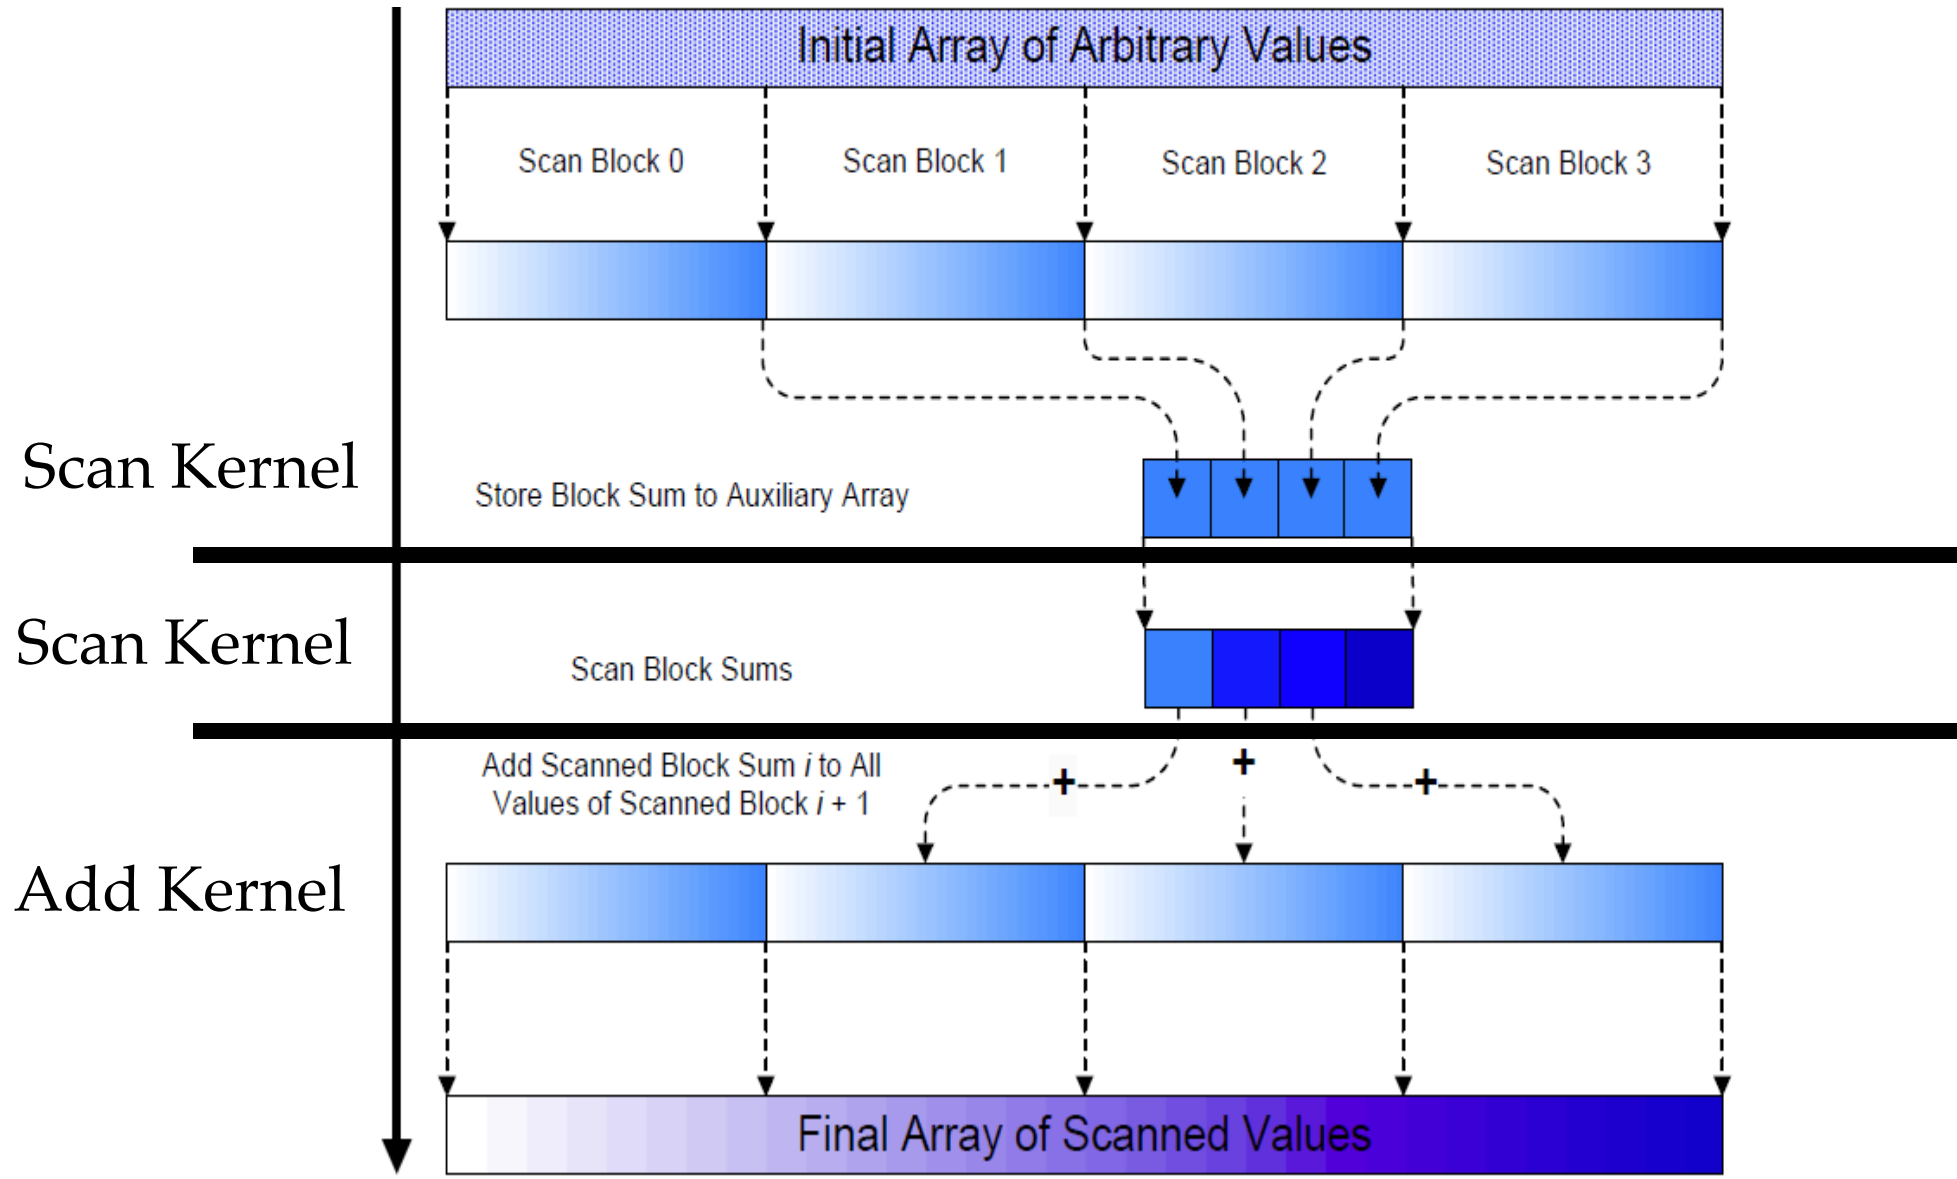
# Scan of Arbitrary Length Input

- Build on the scan kernel that handles up to  $2 \times \text{blockDim.x}$  elements from Brent-Kung
  - For Kogge-Stone, have each section of  $\text{blockDim.x}$  elements assigned to a block
- Have each block write the sum of its section into a Sum array using its  $\text{blockIdx.x}$  as index
- Run parallel scan on the Sum array
  - May need to break down Sum into multiple sections if it is too big for a block
- Add the scanned Sum array values to the elements of corresponding sections



# Overall Flow of Complete Scan

## A Hierarchical Approach



# Exclusive Scan Definition

**Definition:** *The exclusive scan operation takes a binary associative operator  $\oplus$ , and an array of  $n$  elements*

$$[x_0, x_1, \dots, x_{n-1}]$$

*and returns the array*

$$[0, x_0, (x_0 \oplus x_1), \dots, (x_0 \oplus x_1 \oplus \dots \oplus x_{n-2})].$$

**Example:** If  $\oplus$  is addition, then the exclusive scan operation

on  $[3 \ 1 \ 7 \ 0 \ 4 \ 1 \ 6 \ 3],$

would return  $[0 \ 3 \ 4 \ 11 \ 11 \ 15 \ 16 \ 22].$

# Why Exclusive Scan

- To find the beginning address of allocated buffers
- Inclusive and Exclusive scans can be easily derived from each other; it is a matter of convenience

[3 1 7 0 4 1 6 3]

Exclusive [0 3 4 11 11 15 16 22]

Inclusive [3 4 11 11 15 16 22 25]

# A simple exclusive scan kernel

- Adapt an inclusive, Kogge-Stone scan kernel
  - Block 0:
    - Thread 0 loads 0 into (shared)  $XY[0]$
    - Other threads load (global)  $X[threadIdx.x-1]$  into  $XY[threadIdx.x]$
  - All other blocks:
    - All thread load  $X[blockIdx.x*blockDim.x+threadIdx.x-1]$  into  $XY[threadIdx.x]$
- Similar adaption for Brent-Kung kernel but pay attention that each thread loads two elements
  - Only one zero should be loaded
  - All elements should be shifted by only one position

Two vertical lines, one blue and one orange, are positioned on the left side of the slide.

**ANY MORE QUESTIONS?  
READ CHAPTER 8**